

From: Johanna Kougbeadjo
Course: Enterprise security (CMAP 685 9080)
Date: 11/03/2025

Microsoft Dynamics 365 Sales AI Implementation Plan

Introduction

This report presents a comprehensive implementation plan for the secure deployment of Microsoft Dynamics 365 Sales AI within the company's environment. The initiative's objective is twofold: to enhance sales performance through AI-driven insights and to ensure that its deployment complies with applicable cybersecurity and data protection frameworks.

The implementation follows the NIST Artificial Intelligence Risk Management Framework (AI RMF 1.0) (NIST, 2023), which integrates governance, accountability, and risk mitigation principles tailored to AI systems. Each project phase has been aligned with practical tools and controls that will ensure that the solution is deployed securely.

The purpose of this plan is to provide a roadmap for secure implementation as well as an assessment of the challenges and residual risks associated with the deployment of an AI tool in our hybrid cloud environment. The report's structure mirrors the project's governance and operational design: establishing oversight, managing intake, defining decision gates, implementing controls, and continuously assessing outcomes.

1. Governance Structure and Decision-Making

This governance framework established for the implementation of Microsoft Dynamics 365 Sales AI will act as a scalable blueprint for all future AI initiatives within the organization. Designed to be multi-layered and adaptive, it reflects the complexity of AI ecosystems.

This structure integrates traditional cybersecurity governance with AI-specific oversight mechanisms, ensuring that ethical, legal, and operational dimensions are managed cohesively across the entire AI lifecycle. Its objective is to maintain transparency, accountability, and traceability in every decision, from data acquisition to post-deployment monitoring.

Beyond compliance, the governance model functions as a strategic risk-control mechanism aligned with the NIST AI RMF 1.0 GOVERN outcomes (1.1–1.7). It ensures that accountability is clearly distributed, decisions are timely and evidence-based, and the project remains aligned with the organization's defined risk appetite.

In this framework, governance roles are explicitly defined to eliminate overlaps and prevent decision bottlenecks. Each participant (executive sponsor, governance committee, project manager, cybersecurity and data science leads, and compliance officer) has distinct authority boundaries and measurable responsibilities.

The chart below details the composition of the AI Governance Board. It outlines decision ownership and escalation paths across all implementation phases.

Role	Responsibilities and Decision Authority
Executive Sponsor (CIO)	Owns overall project success; approves strategic direction, funding, and full-deployment authorization.
AI Governance Committee	Defines ethical AI principles, ensures regulatory alignment, and validates risk treatment decisions.
Project Manager (Enterprise Security)	Coordinates deliverables, resource allocation, and cross-departmental execution.
Cybersecurity Lead	Implements Zero Trust architecture and continuous monitoring using Microsoft Defender for Cloud. Approves go/no-go at technical checkpoints.
Data Science Lead	Ensures AI model transparency, robustness, and explainability through testing and retraining cycles.
Compliance Officer	Oversees GDPR, CCPA, and EU AI Act adherence; leads DPIAs and manages documentation.
Sales Operations Lead	Validates that AI-driven outputs improve user workflows and support business KPIs.

Decision Gates and Control Points

The project lifecycle integrates four decision gates, each designed to confirm that no major investment or technical progression occurs without validated compliance, performance, and security assurance.

- 1. **Gate 1 – Feasibility and Risk Validation:**
The AI Governance Committee assesses technical feasibility, data readiness, and AI-specific risks such as model bias and explainability gaps.

2. **Gate 2 – Funding Approval:**

Executive management reviews cost-benefit analysis, ensuring alignment with ROI projections and risk appetite.

3. **Gate 3 – Pilot Evaluation:**

The steering group validates model performance against baseline metrics (e.g., 95% prediction accuracy, no critical vulnerabilities).

4. **Gate 4 – Full Implementation:**

The CIO and Compliance Officer jointly approve the production rollout after passing external audit and security certification.

This governance design ensures that the deployment maintains traceability, accountability, and explainability, which are three pillars required by the EU AI Act (ISACA, 2024). Unlike a static committee model, this approach embeds dynamic decision gates that balance innovation velocity with security rigor.

2. Intake Process for AI Use Cases

The intake process must be disciplined and transparent to ensure that all AI initiatives are selected and implemented based on quantifiable business value. This process establishes a standardized, repeatable approach to evaluate, prioritize, and authorize AI use cases across departments while ensuring alignment with corporate strategy, data governance policies, and regulatory requirements.

For this project, the intake pipeline leverages Microsoft Purview as a centralized platform for proposal management and documentation (Kaushik et al., 2023). Business units are required to submit structured proposals detailing the project's objectives, intended outcomes, data sources, projected ROI, and potential risks.

Once submitted, each proposal undergoes a multi-phase evaluation workflow designed to integrate technical, ethical, and operational considerations:

- **Screening Phase:** Confirms data availability, quality, and classification. Proposals that involve restricted or sensitive datasets are immediately flagged for additional privacy impact assessment.
- **Evaluation Phase:** The Enterprise Security and Compliance teams assess each use case against the NIST AI RMF “Map” and “Manage” functions, focusing on data integrity, algorithmic transparency, model robustness, and cybersecurity posture.
- **Prioritization Phase:** A weighted scoring model ranks projects based on business impact (40%), risk level (35%), and technical feasibility (25%), ensuring that high-impact, low-risk initiatives are prioritized.

- **Approval Phase:** The AI Governance Committee reviews top-ranking proposals and authorizes only those that demonstrate compliance readiness, measurable ROI, and alignment with the company's ethical AI principles.

This structured intake process mitigates the risk of “AI sprawl” (the uncontrolled proliferation of isolated tools that fragment oversight) and institutionalizes traceable decision-making. Every evaluation step is logged within Purview, to ensure auditability, accountability, and documentation integrity consistent with the NIST AI RMF, and the EU AI Act.

3. Implementation Plan

Purpose and Objectives

The primary goal of this initiative is to integrate AI-driven forecasting and customer insights into the sales lifecycle while preserving a robust cybersecurity and compliance posture. The deployment of Microsoft Dynamics 365 Sales AI aims to optimize decision-making, streamline CRM operations, and reinforce trust in AI-enabled business intelligence.

The project pursues the following measurable objectives:

- Improve forecasting accuracy by 25% through predictive analytics and automated data correlation.
- Reduce manual CRM data handling by 20% via automation and workflow optimization.
- Achieve 100% compliance with applicable privacy and AI governance frameworks, including the GDPR, CCPA, and EU AI Act.

These targets ensure that the initiative delivers quantifiable business value while aligning with the organization's overarching strategy for secure digital transformation.

Scope and Constraints

The deployment spans global CRM operations, integrating cloud-based Dynamics 365 with on-premise data connectors. This hybrid architecture introduces inherent challenges: multi-layered API dependencies, inconsistent security baselines, and regulatory divergence across jurisdictions (Microsoft, 2022).

To mitigate these challenges, the architecture will leverage Microsoft Defender for Cloud for posture management, Microsoft Entra ID for secure identity governance, and Wiz for continuous visibility and compliance validation across hybrid assets (Wiz.io, 2024). The design follows a Zero Trust architecture, which ensures that authentication, authorization, and verification processes are enforced at every interaction point.

Automated posture management, compliance scanning, and benchmark validation following Microsoft's Cloud Security Benchmark v1 best practices (Microsoft, 2024) will sustain long-term system resilience and reduce administrative overhead.

Strategic Approach

The implementation strategy aligns directly with the AI RMF 1.0 to ensure security, accountability, and reliability throughout the AI lifecycle. The framework's four functions (Govern, Map, Measure, and Manage) provide the structure for both risk management and operational execution:

- **Govern:** Establish a cross-functional AI governance structure with clearly defined roles, reporting mechanisms, and decision-making authority.
- **Map:** Identify, classify, and trace all data flows using Microsoft Purview to ensure transparency, lineage tracking, and privacy compliance.
- **Measure:** Validate AI model accuracy, robustness, and fairness through Protect AI Guardian, Snyk, and the Adversarial Robustness Toolbox (ART).
- **Manage:** Implement continuous monitoring via Microsoft Sentinel, conduct quarterly governance reviews, and document ongoing compliance results.

This structured integration reinforces consistency across data governance, cybersecurity, and AI assurance, cultivating a secure-by-design innovation culture and minimizing the risk of unmonitored AI drift or bias over time.

Execution Phases

The project will be executed over eight months, divided into sequential and interdependent milestones to ensure governance alignment and controlled progress:

1. **Governance Setup and Role Assignment:** Formalize the AI Governance Board, clarify decision-making authority, and validate compliance accountability.
2. **Security Baseline Configuration:** Deploy baseline controls within Defender for Cloud and Entra ID, establishing access management, encryption, and policy enforcement.
3. **Data Classification and DPIA Completion:** Utilize Microsoft Purview to classify sensitive data and finalize Data Protection Impact Assessments (DPIAs).
4. **Model Testing and Validation:** Perform integrity, robustness, and vulnerability testing using Snyk, Guardian, and ART (LFAD, n.d.; Snyk, 2024).
5. **Secure Deployment and Compliance Scanning:** Execute deployment with Wiz to detect drift and maintain compliance against security benchmarks.
6. **Integration with Microsoft Sentinel:** Centralize threat detection, incident correlation, and automated response.

7. **Audit, Documentation, and Training** – Conduct internal audits, update operational documentation, and provide staff training to sustain adoption and governance maturity.

The total budget is estimated at \$450,000. It covers software licensing, personnel, training, and third-party audits. This investment aligns with industry standards for secure AI implementations of comparable scale (AgenticDream, 2025).

Deployment Summary

Purpose of Deployment	Goals (Business & Project Specific)	Scope	Expected Constraints & Barriers	Strategy	Architecture	Step-by-Step Actions	Resources Needed	Estimated Time & Costs
Secure and compliant rollout of Microsoft Dynamics 365 Sales AI	- Increase forecasting accuracy 25% - Reduce manual CRM effort 20% - Maintain 100% GDPR/AI Act compliance	Global hybrid Azure + on-prem CRM environments	Complex API integrations; evolving EU AI Act; data quality variance	Apply NIST AI RMF lifecycle with Zero Trust and continuous governance	Hybrid architecture linking Dynamics 365, Azure, Purview, Defender for Cloud, Wiz, Sentinel	1. Governance setup 2. Baseline security 3. Data mapping via Purview 4. Model validation 5. Secure deployment 6. Monitoring 7. Quarterly reviews	Project Manager; 2 Security Engineers; Data Scientist; Compliance Officer	8 months ≈ \$450,000 (licensing, training, audits)

4. Success Criteria and Performance Measurement

The effectiveness of the Microsoft Dynamics 365 Sales AI implementation will be assessed through a combination of quantitative performance metrics and qualitative governance indicators. Success is defined not only by technical deployment but also by demonstrable business impact, regulatory conformity, and sustained reduction of cyber risk exposure.

Performance measurement will be embedded within the organization’s AI governance cycle, ensuring that metrics are reviewed quarterly and used to drive continuous improvement. The selected criteria focus on operational efficiency, forecasting accuracy, regulatory compliance, security posture, and incident response performance.

Performance Dimension	Metric	Target Value	Measurement Method
Efficiency	Reduction in manual CRM workload	20% decrease in time spent on manual data entry and report generation	Analysis of CRM workflow and automation analytics
Accuracy	Forecast reliability improvement	25% increase in predictive accuracy over six months	Comparative model-performance evaluations using validation datasets
Compliance	Conformance with GDPR, CCPA, and EU AI Act	100% adherence across all operational regions	Microsoft Purview compliance dashboards and quarterly audit reports
Security	Reduction in configuration drift and unpatched vulnerabilities	40% reduction compared with baseline	Continuous posture assessment via Wiz and Defender for Cloud
Response	Mean Time to Detect (MTTD)	≤ 15 minutes per incident	Microsoft Sentinel incident log analysis and SOC response metrics

These metrics provide objective evidence of project success and directly align with the NIST AI RMF “Measure” and “Manage” functions.

The measurement framework ensures that governance teams can validate whether AI outputs deliver consistent, bias-free results; track regulatory compliance dynamically

rather than reactively; and quantify the cybersecurity improvements attributable to the new AI architecture.

Quarterly performance reviews will be conducted by the AI Governance Committee, and any deviation from target thresholds will trigger corrective actions such as model retraining, configuration updates, or process refinement. This approach embeds continuous assurance into the AI lifecycle, which ensures that performance remains aligned with both organizational objectives and risk-tolerance boundaries.

5. Bottom-Line Benefits and Residual Cyber Risk

The deployment of Microsoft Dynamics 365 Sales AI is projected to yield substantial benefits across the organization's operational, compliance, and security domains. From an operational standpoint, the automation of data-intensive CRM activities will release valuable analyst time, allowing teams to focus on higher-value strategic functions such as customer engagement optimization and revenue analysis. These efficiency gains are expected to translate into measurable cost savings and improved productivity across global business units.

In terms of compliance, the integration of Microsoft Purview and Compliance Manager introduces a continuous monitoring and documentation framework that significantly reduces the burden of manual audits. This supports sustained alignment with the GDPR, CCPA, and EU AI Act; and strengthens the company's posture in the face of data protection regulations. The result is a governance environment where regulatory adherence becomes a dynamic and traceable process rather than a reactive effort.

From a security perspective, the combined use of Microsoft Defender for Cloud and Microsoft Sentinel provides real-time visibility into hybrid environments and proactive incident detection capabilities (SANS Institute, 2023). This unified security posture management reduces the organization's exposure to misconfigurations, unauthorized access, and unpatched vulnerabilities, ensuring a consistent defensive baseline across all AI and CRM assets.

Beyond operational and technical gains, the deployment also delivers strategic value. By leveraging predictive analytics, the system enhances the precision of sales forecasting and customer behavior modeling, enabling data-driven decision-making at the executive level. These insights directly contribute to the improvement of customer satisfaction and more reliable revenue projections. This reinforces the strategic role of AI as a business enabler rather than a technical accessory.

Despite these advantages, certain residual risks remain inherent to AI systems. Model drift, data leakage through third-party connectors, and the ongoing evolution of international AI regulations represent persistent areas of concern. These risks are actively mitigated through a combination of regular model retraining, adversarial testing using the Adversarial Robustness Toolbox (ART), and quarterly encryption and vendor security assessments.

Additionally, Purview Compliance Manager is used to monitor continuous regulatory alignment, ensuring that any emerging compliance gaps are addressed promptly.

Overall, the residual risk level is assessed as moderate and acceptable, consistent with the organization's established risk appetite. However, it is recommended that an independent third-party audit be conducted annually to verify system resilience, validate compliance, and confirm that the implemented controls remain effective against evolving AI-specific threats.

This balanced approach, which combines measurable benefits with proactive risk governance, ensures that the AI deployment delivers tangible business value without compromising ethical, regulatory, or security standards.

6. Critical Reflection and Continuous Improvement

A defining strength of this implementation plan lies in its commitment to continuous governance maturity rather than static compliance. This approach recognizes that AI governance must remain iterative, learning from operational outcomes and environmental changes. By embedding review cycles, performance audits, and risk reassessments into the governance structure, the organization ensures that oversight remains relevant as both technology and external obligations evolve. This transforms governance from a reactive mechanism into a strategic enabler of resilience and trust.

Nonetheless, certain limitations persist within the current state of AI security tooling. Advanced attack vectors, such as data poisoning, model inversion, and adversarial manipulation, remain only partially detectable by existing commercial and open-source systems. Recent analysis highlights that open-source dependency vulnerabilities have sharply increased, which amplifies AI model supply-chain risk (Snyk, 2024). To mitigate this limitation, the organization should cultivate strategic partnerships with both technology vendors and open-source communities, such as the Linux Foundation AI & Data initiative (LFAD), to integrate emerging defensive methodologies and threat intelligence into its security operations.

Equally critical to sustained success is the human element of AI governance. Continuous staff training and awareness programs should focus on developing advanced competencies in AI-specific threat modeling, ethical risk assessment, and model interpretability. This investment will ensure that human oversight complements automated monitoring, which will foster a governance environment that balances technological precision with ethical judgment.

Ultimately, the organization's ability to maintain trustworthy and secure AI operations will depend on its willingness to view governance as a living system: one that learns, adapts, and strengthens over time through collaboration, foresight, and continuous improvement.

Conclusion and Recommendation

The secure implementation of Microsoft Dynamics 365 Sales AI demonstrates a strong alignment with the organization's overarching objectives for digital transformation, operational efficiency, and compliance excellence. By embedding Microsoft's advanced security ecosystem within a NIST AI RMF-aligned governance model, the initiative balances innovation with rigorous accountability. This integration ensures that technological progress is matched by ethical oversight, resilience, and long-term sustainability.

While residual risks such as model drift and evolving regulatory requirements cannot be entirely eliminated, they are maintained within acceptable thresholds through continuous monitoring, periodic model retraining, and independent external audits. The governance framework established for this project provides the agility required to adapt to new regulations, emerging threats, and shifting business priorities without compromising trust or security.

The results of this assessment confirm that the project operates well within the organization's defined risk appetite and supports measurable improvements in productivity, decision accuracy, and cybersecurity posture. The AI deployment not only enhances sales intelligence and customer engagement but also strengthens the company's reputation as a responsible and compliant innovator.

Final Recommendation:

Proceed with full implementation following successful pilot validation, supported by an annual third-party compliance review to ensure ongoing alignment with the NIST AI RMF.

References

AgenticDream. (2025). *AI implementation cost guide 2025: Complete breakdown*.

<https://www.agenticdream.com/ai-implementation-cost-guide>

ISACA. (2024). Understanding the EU AI Act: White paper.

<https://www.isaca.org/resources/white-papers/2024/understanding-the-eu-ai-act>

Kaushik, R., Narayan, A., & Zhang, J. (2023). Microsoft Purview: A system for central governance of data. *Proceedings of the VLDB Endowment*, 16(13), 3624–3636.

<https://www.vldb.org/pvldb/vol16/p3624-kaushik.pdf>

Linux Foundation AI & Data (LFAD). (n.d.). Adversarial Robustness Toolbox documentation.

<https://adversarial-robustness-toolbox.readthedocs.io/>

Microsoft. (2022). Microsoft Cloud Security for Enterprise Architects.

<https://go.microsoft.com/fwlink/p/?linkid=842070>

Microsoft. (2024). Microsoft Cloud Security Benchmark v1. <https://learn.microsoft.com/en-us/security/benchmark/azure/overview>

National Institute of Standards and Technology (NIST). (2023). Artificial Intelligence Risk Management Framework (AI RMF 1.0). U.S. Department of Commerce. <https://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf>

SANS Institute. (2023). Securing the future with Microsoft Defender for Cloud: Best practices and insights. <https://www.sans.org/white-papers/>

Snyk. (2024). 2024 State of Open Source Security Report: Slowing progress and new challenges. <https://snyk.io/blog/2024-open-source-security-report-slowing-progress-and-new-challenges-for/>

Wiz.io. (2024). Wiz for Cloud Security Posture Management (Technical Whitepaper). https://www.datocms-assets.com/75231/1686070668-wiz_cspm_technical_whitepaper_final.pdf